

untitled64

November 29, 2024

```
[11]: filename="https://cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud/
↳IBMDeveloperSkillsNetwork-DA0101EN-SkillsNetwork/labs/Data%20files/auto.csv"
headers = ["symboling", "normalized-losses", "make", "fuel-type", "aspiration",
↳"num-of-doors", "body-style",
↳"drive-wheels", "engine-location", "wheel-base",
↳"length", "width", "height", "curb-weight", "engine-type",
↳"num-of-cylinders",
↳"engine-size", "fuel-system", "bore", "stroke", "compression-ratio", "horsepower",
↳"peak-rpm", "city-mpg", "highway-mpg", "price"]
```

```
[13]: import pandas as pd
import matplotlib.pyplot as plt
df = pd.read_csv(filename, names = headers)
```

```
[15]: df
```

```
[15]:      symboling  normalized-losses      make fuel-type aspiration \
0           3           ?  alfa-romero    gas      std
1           3           ?  alfa-romero    gas      std
2           1           ?  alfa-romero    gas      std
3           2          164      audi    gas      std
4           2          164      audi    gas      std
..      ...      ...      ...      ...      ...
200        -1          95     volvo    gas      std
201        -1          95     volvo    gas     turbo
202        -1          95     volvo    gas      std
203        -1          95     volvo  diesel     turbo
204        -1          95     volvo    gas     turbo

      num-of-doors  body-style  drive-wheels  engine-location  wheel-base  ... \
0           two  convertible      rwd      front      88.6  ...
1           two  convertible      rwd      front      88.6  ...
2           two   hatchback      rwd      front      94.5  ...
3           four     sedan      fwd      front      99.8  ...
4           four     sedan      4wd      front      99.4  ...
..      ...      ...      ...      ...      ...
200        four     sedan      rwd      front      109.1  ...
```

201	four	sedan	rwd	front	109.1	...
202	four	sedan	rwd	front	109.1	...
203	four	sedan	rwd	front	109.1	...
204	four	sedan	rwd	front	109.1	...

	engine-size	fuel-system	bore	stroke	compression-ratio	horsepower	\
0	130	mpfi	3.47	2.68	9.0	111	
1	130	mpfi	3.47	2.68	9.0	111	
2	152	mpfi	2.68	3.47	9.0	154	
3	109	mpfi	3.19	3.40	10.0	102	
4	136	mpfi	3.19	3.40	8.0	115	
..	
200	141	mpfi	3.78	3.15	9.5	114	
201	141	mpfi	3.78	3.15	8.7	160	
202	173	mpfi	3.58	2.87	8.8	134	
203	145	idi	3.01	3.40	23.0	106	
204	141	mpfi	3.78	3.15	9.5	114	

	peak-rpm	city-mpg	highway-mpg	price
0	5000	21	27	13495
1	5000	21	27	16500
2	5000	19	26	16500
3	5500	24	30	13950
4	5500	18	22	17450
..
200	5400	23	28	16845
201	5300	19	25	19045
202	5500	18	23	21485
203	4800	26	27	22470
204	5400	19	25	22625

[205 rows x 26 columns]

```
[17]: import numpy as np
df.replace("?", np.nan, inplace=True)
```

```
[19]: df.head(50)
```

```
[19]:      symboling  normalized-losses      make fuel-type aspiration \
0           3           NaN  alfa-romero      gas      std
1           3           NaN  alfa-romero      gas      std
2           1           NaN  alfa-romero      gas      std
3           2          164      audi      gas      std
4           2          164      audi      gas      std
5           2           NaN      audi      gas      std
6           1          158      audi      gas      std
7           1           NaN      audi      gas      std
```

8	1	158	audi	gas	turbo
9	0	NaN	audi	gas	turbo
10	2	192	bmw	gas	std
11	0	192	bmw	gas	std
12	0	188	bmw	gas	std
13	0	188	bmw	gas	std
14	1	NaN	bmw	gas	std
15	0	NaN	bmw	gas	std
16	0	NaN	bmw	gas	std
17	0	NaN	bmw	gas	std
18	2	121	chevrolet	gas	std
19	1	98	chevrolet	gas	std
20	0	81	chevrolet	gas	std
21	1	118	dodge	gas	std
22	1	118	dodge	gas	std
23	1	118	dodge	gas	turbo
24	1	148	dodge	gas	std
25	1	148	dodge	gas	std
26	1	148	dodge	gas	std
27	1	148	dodge	gas	turbo
28	-1	110	dodge	gas	std
29	3	145	dodge	gas	turbo
30	2	137	honda	gas	std
31	2	137	honda	gas	std
32	1	101	honda	gas	std
33	1	101	honda	gas	std
34	1	101	honda	gas	std
35	0	110	honda	gas	std
36	0	78	honda	gas	std
37	0	106	honda	gas	std
38	0	106	honda	gas	std
39	0	85	honda	gas	std
40	0	85	honda	gas	std
41	0	85	honda	gas	std
42	1	107	honda	gas	std
43	0	NaN	isuzu	gas	std
44	1	NaN	isuzu	gas	std
45	0	NaN	isuzu	gas	std
46	2	NaN	isuzu	gas	std
47	0	145	jaguar	gas	std
48	0	NaN	jaguar	gas	std
49	0	NaN	jaguar	gas	std

	num-of-doors	body-style	drive-wheels	engine-location	wheel-base	...	\
0	two	convertible	rwd	front	88.6	...	
1	two	convertible	rwd	front	88.6	...	
2	two	hatchback	rwd	front	94.5	...	

3	four	sedan	fwd	front	99.8	...
4	four	sedan	4wd	front	99.4	...
5	two	sedan	fwd	front	99.8	...
6	four	sedan	fwd	front	105.8	...
7	four	wagon	fwd	front	105.8	...
8	four	sedan	fwd	front	105.8	...
9	two	hatchback	4wd	front	99.5	...
10	two	sedan	rwd	front	101.2	...
11	four	sedan	rwd	front	101.2	...
12	two	sedan	rwd	front	101.2	...
13	four	sedan	rwd	front	101.2	...
14	four	sedan	rwd	front	103.5	...
15	four	sedan	rwd	front	103.5	...
16	two	sedan	rwd	front	103.5	...
17	four	sedan	rwd	front	110.0	...
18	two	hatchback	fwd	front	88.4	...
19	two	hatchback	fwd	front	94.5	...
20	four	sedan	fwd	front	94.5	...
21	two	hatchback	fwd	front	93.7	...
22	two	hatchback	fwd	front	93.7	...
23	two	hatchback	fwd	front	93.7	...
24	four	hatchback	fwd	front	93.7	...
25	four	sedan	fwd	front	93.7	...
26	four	sedan	fwd	front	93.7	...
27	NaN	sedan	fwd	front	93.7	...
28	four	wagon	fwd	front	103.3	...
29	two	hatchback	fwd	front	95.9	...
30	two	hatchback	fwd	front	86.6	...
31	two	hatchback	fwd	front	86.6	...
32	two	hatchback	fwd	front	93.7	...
33	two	hatchback	fwd	front	93.7	...
34	two	hatchback	fwd	front	93.7	...
35	four	sedan	fwd	front	96.5	...
36	four	wagon	fwd	front	96.5	...
37	two	hatchback	fwd	front	96.5	...
38	two	hatchback	fwd	front	96.5	...
39	four	sedan	fwd	front	96.5	...
40	four	sedan	fwd	front	96.5	...
41	four	sedan	fwd	front	96.5	...
42	two	sedan	fwd	front	96.5	...
43	four	sedan	rwd	front	94.3	...
44	two	sedan	fwd	front	94.5	...
45	four	sedan	fwd	front	94.5	...
46	two	hatchback	rwd	front	96.0	...
47	four	sedan	rwd	front	113.0	...
48	four	sedan	rwd	front	113.0	...
49	two	sedan	rwd	front	102.0	...

	engine-size	fuel-system	bore	stroke	compression-ratio	horsepower	\
0	130	mpfi	3.47	2.68	9.00	111	
1	130	mpfi	3.47	2.68	9.00	111	
2	152	mpfi	2.68	3.47	9.00	154	
3	109	mpfi	3.19	3.40	10.00	102	
4	136	mpfi	3.19	3.40	8.00	115	
5	136	mpfi	3.19	3.40	8.50	110	
6	136	mpfi	3.19	3.40	8.50	110	
7	136	mpfi	3.19	3.40	8.50	110	
8	131	mpfi	3.13	3.40	8.30	140	
9	131	mpfi	3.13	3.40	7.00	160	
10	108	mpfi	3.50	2.80	8.80	101	
11	108	mpfi	3.50	2.80	8.80	101	
12	164	mpfi	3.31	3.19	9.00	121	
13	164	mpfi	3.31	3.19	9.00	121	
14	164	mpfi	3.31	3.19	9.00	121	
15	209	mpfi	3.62	3.39	8.00	182	
16	209	mpfi	3.62	3.39	8.00	182	
17	209	mpfi	3.62	3.39	8.00	182	
18	61	2bbl	2.91	3.03	9.50	48	
19	90	2bbl	3.03	3.11	9.60	70	
20	90	2bbl	3.03	3.11	9.60	70	
21	90	2bbl	2.97	3.23	9.41	68	
22	90	2bbl	2.97	3.23	9.40	68	
23	98	mpfi	3.03	3.39	7.60	102	
24	90	2bbl	2.97	3.23	9.40	68	
25	90	2bbl	2.97	3.23	9.40	68	
26	90	2bbl	2.97	3.23	9.40	68	
27	98	mpfi	3.03	3.39	7.60	102	
28	122	2bbl	3.34	3.46	8.50	88	
29	156	mfi	3.60	3.90	7.00	145	
30	92	1bbl	2.91	3.41	9.60	58	
31	92	1bbl	2.91	3.41	9.20	76	
32	79	1bbl	2.91	3.07	10.10	60	
33	92	1bbl	2.91	3.41	9.20	76	
34	92	1bbl	2.91	3.41	9.20	76	
35	92	1bbl	2.91	3.41	9.20	76	
36	92	1bbl	2.92	3.41	9.20	76	
37	110	1bbl	3.15	3.58	9.00	86	
38	110	1bbl	3.15	3.58	9.00	86	
39	110	1bbl	3.15	3.58	9.00	86	
40	110	1bbl	3.15	3.58	9.00	86	
41	110	mpfi	3.15	3.58	9.00	101	
42	110	2bbl	3.15	3.58	9.10	100	
43	111	2bbl	3.31	3.23	8.50	78	
44	90	2bbl	3.03	3.11	9.60	70	

45	90	2bbl	3.03	3.11	9.60	70
46	119	spfi	3.43	3.23	9.20	90
47	258	mpfi	3.63	4.17	8.10	176
48	258	mpfi	3.63	4.17	8.10	176
49	326	mpfi	3.54	2.76	11.50	262

	peak-rpm	city-mpg	highway-mpg	price
0	5000	21	27	13495
1	5000	21	27	16500
2	5000	19	26	16500
3	5500	24	30	13950
4	5500	18	22	17450
5	5500	19	25	15250
6	5500	19	25	17710
7	5500	19	25	18920
8	5500	17	20	23875
9	5500	16	22	NaN
10	5800	23	29	16430
11	5800	23	29	16925
12	4250	21	28	20970
13	4250	21	28	21105
14	4250	20	25	24565
15	5400	16	22	30760
16	5400	16	22	41315
17	5400	15	20	36880
18	5100	47	53	5151
19	5400	38	43	6295
20	5400	38	43	6575
21	5500	37	41	5572
22	5500	31	38	6377
23	5500	24	30	7957
24	5500	31	38	6229
25	5500	31	38	6692
26	5500	31	38	7609
27	5500	24	30	8558
28	5000	24	30	8921
29	5000	19	24	12964
30	4800	49	54	6479
31	6000	31	38	6855
32	5500	38	42	5399
33	6000	30	34	6529
34	6000	30	34	7129
35	6000	30	34	7295
36	6000	30	34	7295
37	5800	27	33	7895
38	5800	27	33	9095
39	5800	27	33	8845

40	5800	27	33	10295
41	5800	24	28	12945
42	5500	25	31	10345
43	4800	24	29	6785
44	5400	38	43	NaN
45	5400	38	43	NaN
46	5000	24	29	11048
47	4750	15	19	32250
48	4750	15	19	35550
49	5000	13	17	36000

[50 rows x 26 columns]

```
[21]: missing_data = df.isnull()
missing_data.head(5)
```

```
[21]:  symboling  normalized-losses  make  fuel-type  aspiration  num-of-doors  \
0      False                True  False      False      False      False
1      False                True  False      False      False      False
2      False                True  False      False      False      False
3      False                False  False      False      False      False
4      False                False  False      False      False      False

      body-style  drive-wheels  engine-location  wheel-base  ...  engine-size  \
0      False      False      False      False  ...      False
1      False      False      False      False  ...      False
2      False      False      False      False  ...      False
3      False      False      False      False  ...      False
4      False      False      False      False  ...      False

      fuel-system  bore  stroke  compression-ratio  horsepower  peak-rpm  \
0      False  False  False      False      False      False
1      False  False  False      False      False      False
2      False  False  False      False      False      False
3      False  False  False      False      False      False
4      False  False  False      False      False      False

      city-mpg  highway-mpg  price
0      False      False  False
1      False      False  False
2      False      False  False
3      False      False  False
4      False      False  False
```

[5 rows x 26 columns]

```
[23]: for column in missing_data.columns.values.tolist():  
      print(column)  
      print (missing_data[column].value_counts())  
      print("")
```

```
symboling  
symboling  
False    205  
Name: count, dtype: int64
```

```
normalized-losses  
normalized-losses  
False    164  
True      41  
Name: count, dtype: int64
```

```
make  
make  
False    205  
Name: count, dtype: int64
```

```
fuel-type  
fuel-type  
False    205  
Name: count, dtype: int64
```

```
aspiration  
aspiration  
False    205  
Name: count, dtype: int64
```

```
num-of-doors  
num-of-doors  
False    203  
True       2  
Name: count, dtype: int64
```

```
body-style  
body-style  
False    205  
Name: count, dtype: int64
```

```
drive-wheels  
drive-wheels  
False    205  
Name: count, dtype: int64
```


engine-location
engine-location
False 205
Name: count, dtype: int64

wheel-base
wheel-base
False 205
Name: count, dtype: int64

length
length
False 205
Name: count, dtype: int64

width
width
False 205
Name: count, dtype: int64

height
height
False 205
Name: count, dtype: int64

curb-weight
curb-weight
False 205
Name: count, dtype: int64

engine-type
engine-type
False 205
Name: count, dtype: int64

num-of-cylinders
num-of-cylinders
False 205
Name: count, dtype: int64

engine-size
engine-size
False 205
Name: count, dtype: int64

fuel-system
fuel-system
False 205

Name: count, dtype: int64

bore

bore

False 201

True 4

Name: count, dtype: int64

stroke

stroke

False 201

True 4

Name: count, dtype: int64

compression-ratio

compression-ratio

False 205

Name: count, dtype: int64

horsepower

horsepower

False 203

True 2

Name: count, dtype: int64

peak-rpm

peak-rpm

False 203

True 2

Name: count, dtype: int64

city-mpg

city-mpg

False 205

Name: count, dtype: int64

highway-mpg

highway-mpg

False 205

Name: count, dtype: int64

price

price

False 201

True 4

Name: count, dtype: int64

```
[25]: avg_norm_loss = df["normalized-losses"].astype("float").mean(axis=0)
      print("Average of normalized-losses:", avg_norm_loss)
```

Average of normalized-losses: 122.0

```
[27]: df["normalized-losses"].replace(np.nan, avg_norm_loss, inplace=True)
```

C:\Users\USER\AppData\Local\Temp\ipykernel_13824\2599940699.py:1: FutureWarning:
A value is trying to be set on a copy of a DataFrame or Series through chained
assignment using an inplace method.

The behavior will change in pandas 3.0. This inplace method will never work
because the intermediate object on which we are setting values always behaves as
a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using
'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value)
instead, to perform the operation inplace on the original object.

```
df["normalized-losses"].replace(np.nan, avg_norm_loss, inplace=True)
```

```
[29]: avg_bore=df['bore'].astype('float').mean(axis=0)
      print("Average of bore:", avg_bore)
```

Average of bore: 3.3297512437810943

```
[31]: df["bore"].replace(np.nan, avg_bore, inplace=True)
```

C:\Users\USER\AppData\Local\Temp\ipykernel_13824\1952189479.py:1: FutureWarning:
A value is trying to be set on a copy of a DataFrame or Series through chained
assignment using an inplace method.

The behavior will change in pandas 3.0. This inplace method will never work
because the intermediate object on which we are setting values always behaves as
a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using
'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value)
instead, to perform the operation inplace on the original object.

```
df["bore"].replace(np.nan, avg_bore, inplace=True)
```

```
[33]: df["stroke"].replace(np.nan,df['stroke'].astype('float').mean(axis=0))
```

```
[33]: 0      2.68
      1      2.68
      2      3.47
      3      3.40
```

```

4      3.40
...
200    3.15
201    3.15
202    2.87
203    3.40
204    3.15
Name: stroke, Length: 205, dtype: object

```

```
[35]: df["stroke"]
```

```

[35]: 0      2.68
1      2.68
2      3.47
3      3.40
4      3.40
...
200    3.15
201    3.15
202    2.87
203    3.40
204    3.15
Name: stroke, Length: 205, dtype: object

```

```
[37]: df
```

```

[37]:      symboling  normalized-losses      make fuel-type aspiration \
0           3           122.0  alfa-romero    gas      std
1           3           122.0  alfa-romero    gas      std
2           1           122.0  alfa-romero    gas      std
3           2            164      audi      gas      std
4           2            164      audi      gas      std
..      ...      ...      ...      ...      ...
200        -1            95     volvo      gas      std
201        -1            95     volvo      gas    turbo
202        -1            95     volvo      gas      std
203        -1            95     volvo  diesel    turbo
204        -1            95     volvo      gas    turbo

      num-of-doors  body-style  drive-wheels  engine-location  wheel-base  ... \
0           two  convertible      rwd      front      88.6  ...
1           two  convertible      rwd      front      88.6  ...
2           two   hatchback      rwd      front      94.5  ...
3           four      sedan      fwd      front      99.8  ...
4           four      sedan      4wd      front      99.4  ...
..      ...      ...      ...      ...      ...
200        four      sedan      rwd      front      109.1  ...

```

201	four	sedan	rwd	front	109.1	...
202	four	sedan	rwd	front	109.1	...
203	four	sedan	rwd	front	109.1	...
204	four	sedan	rwd	front	109.1	...

	engine-size	fuel-system	bore	stroke	compression-ratio	horsepower	\
0	130	mpfi	3.47	2.68	9.0	111	
1	130	mpfi	3.47	2.68	9.0	111	
2	152	mpfi	2.68	3.47	9.0	154	
3	109	mpfi	3.19	3.40	10.0	102	
4	136	mpfi	3.19	3.40	8.0	115	
..	
200	141	mpfi	3.78	3.15	9.5	114	
201	141	mpfi	3.78	3.15	8.7	160	
202	173	mpfi	3.58	2.87	8.8	134	
203	145	idi	3.01	3.40	23.0	106	
204	141	mpfi	3.78	3.15	9.5	114	

	peak-rpm	city-mpg	highway-mpg	price
0	5000	21	27	13495
1	5000	21	27	16500
2	5000	19	26	16500
3	5500	24	30	13950
4	5500	18	22	17450
..
200	5400	23	28	16845
201	5300	19	25	19045
202	5500	18	23	21485
203	4800	26	27	22470
204	5400	19	25	22625

[205 rows x 26 columns]

```
[39]: df["horsepower"].astype('float').mean(axis=0)
```

```
[39]: 104.25615763546799
```

```
[41]: df['horsepower'].replace(np.nan,df["horsepower"].astype('float').mean(axis=0) ,
    ↪inplace=True)
```

C:\Users\USER\AppData\Local\Temp\ipykernel_13824\3745643968.py:1: FutureWarning:
A value is trying to be set on a copy of a DataFrame or Series through chained
assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never work
because the intermediate object on which we are setting values always behaves as
a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
df['horsepower'].replace(np.nan,df["horsepower"].astype('float').mean(axis=0), inplace=True)
```

```
[43]: avg_peakrpm=df['peak-rpm'].astype('float').mean(axis=0)
      print("Average peak rpm:", avg_peakrpm)
```

Average peak rpm: 5125.369458128079

```
[45]: df['peak-rpm'].replace(np.nan, avg_peakrpm, inplace=True)
```

C:\Users\USER\AppData\Local\Temp\ipykernel_13824\2061375298.py:1: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
df['peak-rpm'].replace(np.nan, avg_peakrpm, inplace=True)
```

```
[47]: df['num-of-doors'].value_counts()
```

```
[47]: num-of-doors
      four      114
      two       89
      Name: count, dtype: int64
```

```
[49]: df['num-of-doors'].value_counts().idxmax()
```

```
[49]: 'four'
```

```
[51]: #replace the missing 'num-of-doors' values by the most frequent
      df["num-of-doors"].replace(np.nan, "four", inplace=True)
```

C:\Users\USER\AppData\Local\Temp\ipykernel_13824\2406474689.py:2: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as

a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
df["num-of-doors"].replace(np.nan, "four", inplace=True)
```

```
[53]: # simply drop whole row with NaN in "price" column
df.dropna(subset=["price"], axis=0, inplace=True)

# reset index, because we dropped two rows
df.reset_index(drop=True, inplace=True)
```

```
[74]: df.head()
```

```
[74]:  symboling normalized-losses      make fuel-type aspiration num-of-doors \
0         3           122.0  alfa-romero    gas         std         two
1         3           122.0  alfa-romero    gas         std         two
2         1           122.0  alfa-romero    gas         std         two
3         2            164      audi      gas         std         four
4         2            164      audi      gas         std         four

      body-style drive-wheels engine-location  wheel-base  ...  engine-size  \
0  convertible         rwd         front      88.6  ...        130
1  convertible         rwd         front      88.6  ...        130
2   hatchback         rwd         front      94.5  ...        152
3         sedan         fwd         front      99.8  ...        109
4         sedan         4wd         front      99.4  ...        136

      fuel-system  bore  stroke  compression-ratio  horsepower  peak-rpm  city-mpg  \
0         mpfi  3.47   2.68           9.0         111      5000        21
1         mpfi  3.47   2.68           9.0         111      5000        21
2         mpfi  2.68   3.47           9.0         154      5000        19
3         mpfi  3.19   3.40          10.0         102      5500        24
4         mpfi  3.19   3.40           8.0         115      5500        18

      highway-mpg  price
0         27  13495
1         27  16500
2         26  16500
3         30  13950
4         22  17450
```

```
[5 rows x 26 columns]
```

```
[55]: df.dtypes
```

```
[55]: symboling          int64
normalized-losses      object
make                   object
fuel-type              object
aspiration             object
num-of-doors           object
body-style             object
drive-wheels           object
engine-location        object
wheel-base            float64
length                float64
width                  float64
height                 float64
curb-weight            int64
engine-type            object
num-of-cylinders       object
engine-size            int64
fuel-system            object
bore                   object
stroke                 object
compression-ratio      float64
horsepower             object
peak-rpm              object
city-mpg               int64
highway-mpg            int64
price                  object
dtype: object
```

```
[57]: df[["bore", "stroke"]] = df[["bore", "stroke"]].astype("float")
df[["normalized-losses"]] = df[["normalized-losses"]].astype("int")
df[["price"]] = df[["price"]].astype("float")
df[["peak-rpm"]] = df[["peak-rpm"]].astype("float")
```

```
[80]: df.dtypes
```

```
[80]: symboling          int64
normalized-losses      int32
make                   object
fuel-type              object
aspiration             object
num-of-doors           object
body-style             object
drive-wheels           object
engine-location        object
wheel-base            float64
```



```

length          float64
width           float64
height          float64
curb-weight     int64
engine-type     object
num-of-cylinders object
engine-size     int64
fuel-system     object
bore            float64
stroke          float64
compression-ratio float64
horsepower      object
peak-rpm        float64
city-mpg        int64
highway-mpg     int64
price           float64
dtype: object

```

```
[82]: df.head()
```

```

[82]:   symboling  normalized-losses      make fuel-type aspiration \
0         3           122  alfa-romero    gas      std
1         3           122  alfa-romero    gas      std
2         1           122  alfa-romero    gas      std
3         2           164      audi    gas      std
4         2           164      audi    gas      std

   num-of-doors  body-style drive-wheels engine-location  wheel-base  ... \
0         two  convertible      rwd      front      88.6  ...
1         two  convertible      rwd      front      88.6  ...
2         two   hatchback      rwd      front      94.5  ...
3         four      sedan      fwd      front      99.8  ...
4         four      sedan      4wd      front      99.4  ...

   engine-size  fuel-system  bore  stroke  compression-ratio  horsepower  \
0         130      mpfi  3.47   2.68           9.0         111
1         130      mpfi  3.47   2.68           9.0         111
2         152      mpfi  2.68   3.47           9.0         154
3         109      mpfi  3.19   3.40          10.0         102
4         136      mpfi  3.19   3.40           8.0         115

   peak-rpm  city-mpg  highway-mpg   price
0    5000.0      21      27  13495.0
1    5000.0      21      27  16500.0
2    5000.0      19      26  16500.0
3    5500.0      24      30  13950.0
4    5500.0      18      22  17450.0

```

[5 rows x 26 columns]

```
[59]: df["highway-mpg"].replace("mpg", "L/100km", inplace=True)
```

C:\Users\USER\AppData\Local\Temp\ipykernel_13824\460410175.py:1: FutureWarning:
A value is trying to be set on a copy of a DataFrame or Series through chained
assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never work
because the intermediate object on which we are setting values always behaves as
a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using
'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value)
instead, to perform the operation inplace on the original object.

```
df["highway-mpg"].replace("mpg", "L/100km", inplace=True)
```

```
[86]: df
```

```
[86]:
```

	symboling	normalized-losses	make	fuel-type	aspiration	\
0	3	122	alfa-romero	gas	std	
1	3	122	alfa-romero	gas	std	
2	1	122	alfa-romero	gas	std	
3	2	164	audi	gas	std	
4	2	164	audi	gas	std	
..	
196	-1	95	volvo	gas	std	
197	-1	95	volvo	gas	turbo	
198	-1	95	volvo	gas	std	
199	-1	95	volvo	diesel	turbo	
200	-1	95	volvo	gas	turbo	

	num-of-doors	body-style	drive-wheels	engine-location	wheel-base	...	\
0	two	convertible	rwd	front	88.6	...	
1	two	convertible	rwd	front	88.6	...	
2	two	hatchback	rwd	front	94.5	...	
3	four	sedan	fwd	front	99.8	...	
4	four	sedan	4wd	front	99.4	...	
..	
196	four	sedan	rwd	front	109.1	...	
197	four	sedan	rwd	front	109.1	...	
198	four	sedan	rwd	front	109.1	...	
199	four	sedan	rwd	front	109.1	...	
200	four	sedan	rwd	front	109.1	...	

	engine-size	fuel-system	bore	stroke	compression-ratio	horsepower	\
0	130	mpfi	3.47	2.68	9.0	111	
1	130	mpfi	3.47	2.68	9.0	111	
2	152	mpfi	2.68	3.47	9.0	154	
3	109	mpfi	3.19	3.40	10.0	102	
4	136	mpfi	3.19	3.40	8.0	115	
..	
196	141	mpfi	3.78	3.15	9.5	114	
197	141	mpfi	3.78	3.15	8.7	160	
198	173	mpfi	3.58	2.87	8.8	134	
199	145	idi	3.01	3.40	23.0	106	
200	141	mpfi	3.78	3.15	9.5	114	

	peak-rpm	city-mpg	highway-mpg	price
0	5000.0	21	27	13495.0
1	5000.0	21	27	16500.0
2	5000.0	19	26	16500.0
3	5500.0	24	30	13950.0
4	5500.0	18	22	17450.0
..
196	5400.0	23	28	16845.0
197	5300.0	19	25	19045.0
198	5500.0	18	23	21485.0
199	4800.0	26	27	22470.0
200	5400.0	19	25	22625.0

[201 rows x 26 columns]

```
[61]: # transform mpg to L/100km by mathematical operation (235 divided by mpg)
df["highway-mpg"] = 235/df["highway-mpg"]

# rename column name from "highway-mpg" to "highway-L/100km"
df.rename(columns={"highway-mpg": "highway-L/100km"}, inplace=True)

# check your transformed data
df.head()
df["height"]
```

```
[61]: 0      48.8
      1      48.8
      2      52.4
      3      54.3
      4      54.3
      ...
     196      55.5
     197      55.5
     198      55.5
```

```

199    55.5
200    55.5
Name: height, Length: 201, dtype: float64

```

```

[63]: # Write your code below and press Shift+Enter to execute
df["height"]=df["height"]/df["height"].max()

```

```

[91]: df

```

```

[91]:      symboling  normalized-losses      make fuel-type aspiration \
0         3         122  alfa-romero      gas      std
1         3         122  alfa-romero      gas      std
2         1         122  alfa-romero      gas      std
3         2         164      audi      gas      std
4         2         164      audi      gas      std
..      ...      ...      ...      ...      ...
196      -1         95     volvo      gas      std
197      -1         95     volvo      gas     turbo
198      -1         95     volvo      gas      std
199      -1         95     volvo  diesel     turbo
200      -1         95     volvo      gas     turbo

      num-of-doors  body-style drive-wheels engine-location  wheel-base  ... \
0         two  convertible      rwd      front      88.6  ...
1         two  convertible      rwd      front      88.6  ...
2         two   hatchback      rwd      front      94.5  ...
3         four     sedan      fwd      front      99.8  ...
4         four     sedan      4wd      front      99.4  ...
..      ...      ...      ...      ...      ...
196     four     sedan      rwd      front      109.1  ...
197     four     sedan      rwd      front      109.1  ...
198     four     sedan      rwd      front      109.1  ...
199     four     sedan      rwd      front      109.1  ...
200     four     sedan      rwd      front      109.1  ...

      engine-size  fuel-system  bore  stroke  compression-ratio  horsepower  \
0         130      mpfi  3.47    2.68          9.0          111
1         130      mpfi  3.47    2.68          9.0          111
2         152      mpfi  2.68    3.47          9.0          154
3         109      mpfi  3.19    3.40         10.0          102
4         136      mpfi  3.19    3.40          8.0          115
..      ...      ...      ...      ...      ...
196         141      mpfi  3.78    3.15          9.5          114
197         141      mpfi  3.78    3.15          8.7          160
198         173      mpfi  3.58    2.87          8.8          134
199         145      idi  3.01    3.40         23.0          106
200         141      mpfi  3.78    3.15          9.5          114

```

	peak-rpm	city-mpg	highway-mpg	price
0	5000.0	21	8.703704	13495.0
1	5000.0	21	8.703704	16500.0
2	5000.0	19	9.038462	16500.0
3	5500.0	24	7.833333	13950.0
4	5500.0	18	10.681818	17450.0
..
196	5400.0	23	8.392857	16845.0
197	5300.0	19	9.400000	19045.0
198	5500.0	18	10.217391	21485.0
199	4800.0	26	8.703704	22470.0
200	5400.0	19	9.400000	22625.0

[201 rows x 26 columns]

```
[65]: df["height"]
```


```
[65]: 0      0.816054
      1      0.816054
      2      0.876254
      3      0.908027
      4      0.908027
      ...
      196    0.928094
      197    0.928094
      198    0.928094
      199    0.928094
      200    0.928094
      Name: height, Length: 201, dtype: float64
```

```
[67]: df["horsepower"]=df["horsepower"].astype(int, copy=True)
```

```
[69]: %matplotlib inline
import matplotlib as plt
from matplotlib import pyplot
plt.pyplot.hist(df["horsepower"])

# set x/y labels and plot title
plt.pyplot.xlabel("horsepower")
plt.pyplot.ylabel("count")
plt.pyplot.title("horsepower bins")
```

```
[69]: Text(0.5, 1.0, 'horsepower bins')
```



```
[71]: bins = np.linspace(min(df["horsepower"]), max(df["horsepower"]), 4)
      bins
```

```
[71]: array([ 48.          , 119.33333333, 190.66666667, 262.          ])
```

```
[73]: group_names = ['Low', 'Medium', 'High']
```

```
[75]: df['horsepower-binned'] = pd.cut(df['horsepower'], bins, labels=group_names,
      ↪include_lowest=True )
      df[['horsepower', 'horsepower-binned']].head(20)
```

```
[75]:
```

	horsepower	horsepower-binned
0	111	Low
1	111	Low
2	154	Medium
3	102	Low
4	115	Low
5	110	Low
6	110	Low
7	110	Low

8	140	Medium
9	101	Low
10	101	Low
11	121	Medium
12	121	Medium
13	121	Medium
14	182	Medium
15	182	Medium
16	182	Medium
17	48	Low
18	70	Low
19	70	Low


```
[77]: df["horsepower-binned"].value_counts()
```

```
[77]: horsepower-binned
Low      153
Medium   43
High      5
Name: count, dtype: int64
```

```
[103]: %matplotlib inline
import matplotlib as plt
from matplotlib import pyplot
pyplot.bar(group_names, df["horsepower-binned"].value_counts())

# set x/y labels and plot title
plt.pyplot.xlabel("horsepower")
plt.pyplot.ylabel("count")
plt.pyplot.title("horsepower bins")
```

```
[103]: Text(0.5, 1.0, 'horsepower bins')
```




```
[79]: %matplotlib inline
import matplotlib as plt
from matplotlib import pyplot

# draw histogram of attribute "horsepower" with bins = 3
plt.pyplot.hist(df["horsepower"], bins = 3)

# set x/y labels and plot title
plt.pyplot.xlabel("horsepower")
plt.pyplot.ylabel("count")
plt.pyplot.title("horsepower bins")
```

```
[79]: Text(0.5, 1.0, 'horsepower bins')
```

```
[105]: df.columns
```

```
[105]: Index(['symboling', 'normalized-losses', 'make', 'fuel-type', 'aspiration',
            'num-of-doors', 'body-style', 'drive-wheels', 'engine-location',
            'wheel-base', 'length', 'width', 'height', 'curb-weight', 'engine-type',
            'num-of-cylinders', 'engine-size', 'fuel-system', 'bore', 'stroke',
            'compression-ratio', 'horsepower', 'peak-rpm', 'city-mpg',
            'highway-mpg', 'price', 'horsepower-binned'],
            dtype='object')
```

```
[81]: dummy_variable_1 = pd.get_dummies(df["fuel-type"])
      dummy_variable_1.head()
```

```
[81]:   diesel  gas
0   False  True
1   False  True
2   False  True
3   False  True
4   False  True
```

```
[83]: dummy_variable_1.rename(columns={'gas': 'fuel-type-gas', 'diesel':
    ↪ 'fuel-type-diesel'}, inplace=True)
dummy_variable_1.head()
```

```
[83]:    fuel-type-diesel  fuel-type-gas
0                False             True
1                False             True
2                False             True
3                False             True
4                False             True
```

```
[85]: # merge data frame "df" and "dummy_variable_1"
df = pd.concat([df, dummy_variable_1], axis=1)

# drop original column "fuel-type" from "df"
df.drop("fuel-type", axis = 1, inplace=True)
```

```
[87]: # get indicator variables of aspiration and assign it to data frame
    ↪ "dummy_variable_2"
dummy_variable_2 = pd.get_dummies(df['aspiration'])

# change column names for clarity
dummy_variable_2.rename(columns={'std': 'aspiration-std', 'turbo':
    ↪ 'aspiration-turbo'}, inplace=True)

# show first 5 instances of data frame "dummy_variable_1"
dummy_variable_2.head()
```

```
[87]:    aspiration-std  aspiration-turbo
0                True             False
1                True             False
2                True             False
3                True             False
4                True             False
```

```
[89]: # merge the new dataframe to the original dataframe
df = pd.concat([df, dummy_variable_2], axis=1)

# drop original column "aspiration" from "df"
df.drop('aspiration', axis = 1, inplace=True)
```

```
[123]: df.to_csv('clean_df.csv')
```

```
[91]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
```